

# Cats and Dogs CNN Classifier

Ho Huyen Trang

November 29, 2025

# Summary

Convolutional neural networks (CNNs) are central to modern computer vision, as they learn hierarchical feature representations directly from images. Beyond traditional supervised learning, recent advances incorporate self-supervision, weak supervision, large-scale pretraining, and knowledge distillation to reduce reliance on labeled data.

In this project, a lightweight CNN (MobileNetV2) is trained for cat–dog classification using two strategies:

- **Self-supervised contrastive learning (SimCLR)**
- **Knowledge distillation** from a pre-trained ResNet50 teacher

Experimental results show that a MobileNetV2 student trained purely via knowledge distillation reaches — and sometimes exceeds — the performance of one initialized with SimCLR pretraining, while requiring significantly less training time.

# Problem Statement

Deep CNNs offer strong performance in visual recognition but typically require:

- large labeled datasets, and
- computationally heavy architectures.

Lightweight models are efficient but can't be as concise trained from scratch on limited data. Knowledge transfer techniques—transfer learning, knowledge distillation, and feature-based adaptation—help smaller models inherit rich representations and improve accuracy.

Large labeled datasets are also costly, motivating self-supervised learning methods such as **SimCLR** that extract features from unlabeled data.

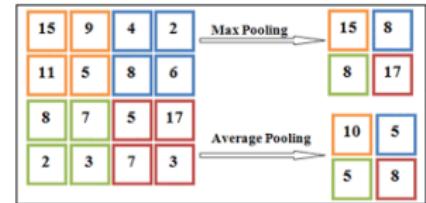
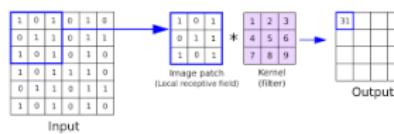
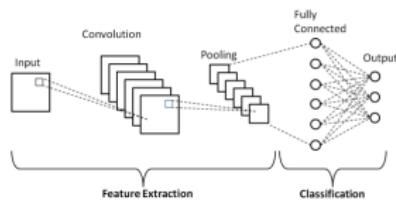
**This project evaluates two techniques—SimCLR and logits-based knowledge distillation—to enhance MobileNetV2 on a cats vs. dogs classification task under limited resources (data + computation).**

# Methods: CNN Structure

**Convolutional Neural Networks (CNNs)** extract hierarchical spatial features from images through:

- **Convolution layers** — learn local patterns (edges, textures)
- **Pooling layers** — reduce spatial resolution and noise
- **Nonlinear activations (ReLU)** — introduce nonlinearity
- **Fully connected / classifier head** — map features to labels

# CNN Structure



Examples of CNN components: full architecture (left), convolution operation (middle), and pooling operation (right).

# Methods: SimCLR (Self-Supervised Learning)

**SimCLR** learns image representations by maximizing agreement between augmented views of the same image.

- Create two augmented views ( $x_i, x_j$ ) of each image.
- Encode both using the same CNN encoder + projection head.
- Use the **NT-Xent contrastive loss**:

$$\mathcal{L}_{i,j} = -\log \frac{\exp(\text{sim}(z_i, z_j)/\tau)}{\sum_{k=1}^{2N} \mathbf{1}_{[k \neq i]} \exp(\text{sim}(z_i, z_k)/\tau)}$$

- $\text{sim}(u, v)$ : cosine similarity
- $\tau$ : temperature
- $2N$ : number of augmented samples in the batch

This pulls positive pairs together and pushes all other samples apart.

## Feature Maps for MobileNetV2 - Layer: features.0.2

Channel 0



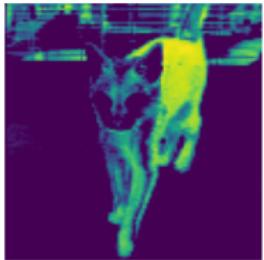
Channel 1



Channel 2



Channel 3



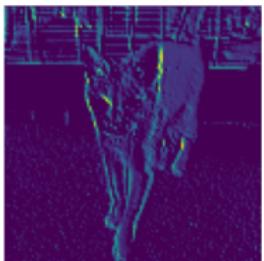
Channel 4



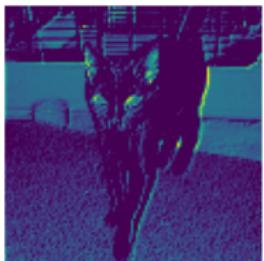
Channel 5



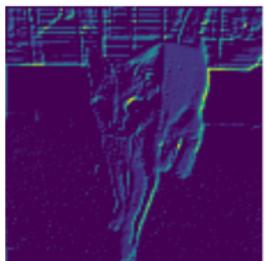
Channel 6



Channel 7



Channel 8



Channel 9



Channel 10



Channel 11



Channel 12

Channel 13

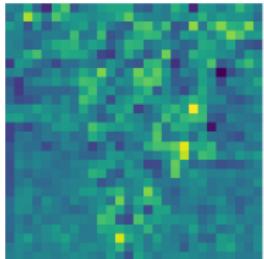
Channel 14

Channel 15

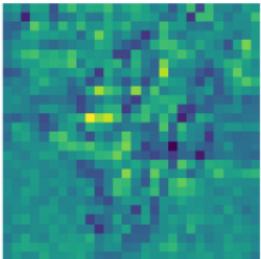


## Feature Maps for MobileNetV2 - Layer: features.6.conv.2

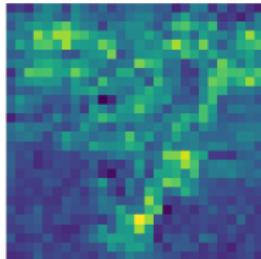
Channel 0



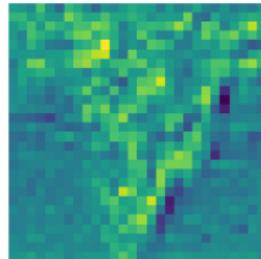
Channel 1



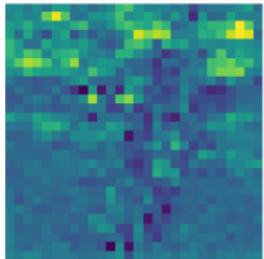
Channel 2



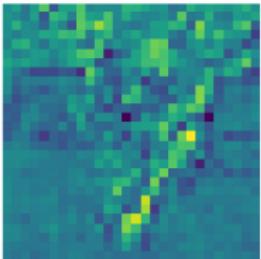
Channel 3



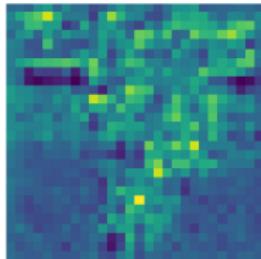
Channel 4



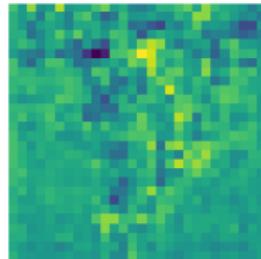
Channel 5



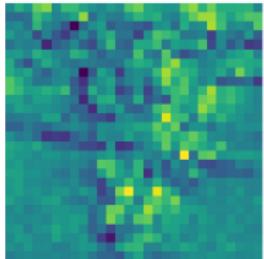
Channel 6



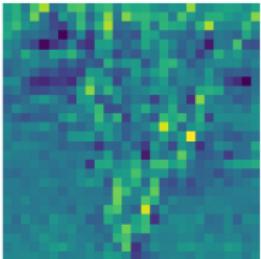
Channel 7



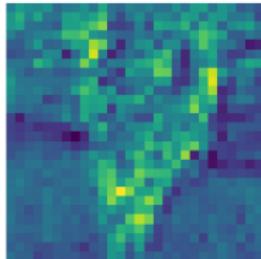
Channel 8



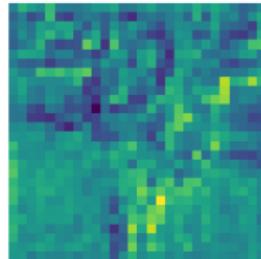
Channel 9



Channel 10



Channel 11



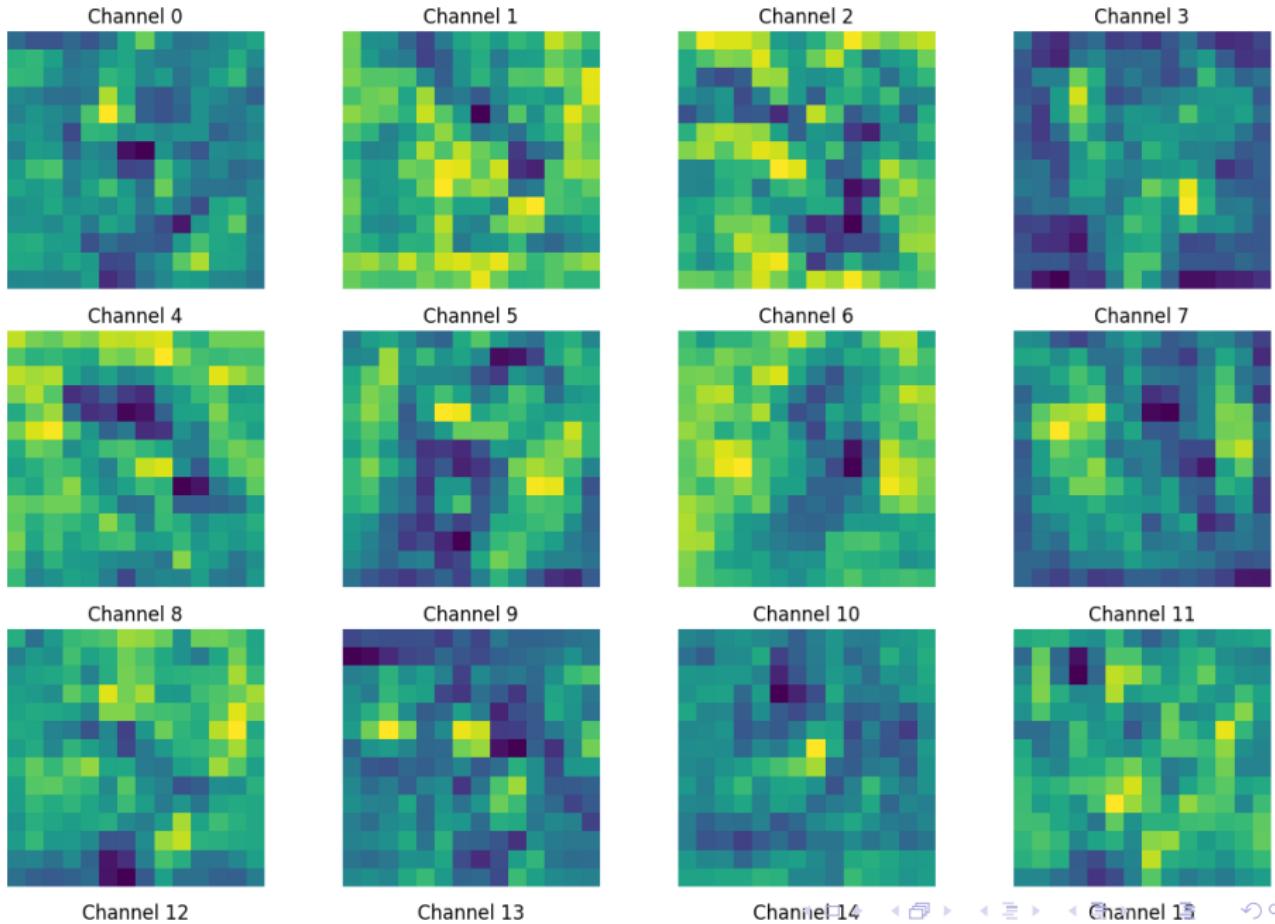
Channel 12

Channel 13

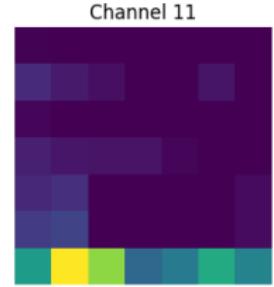
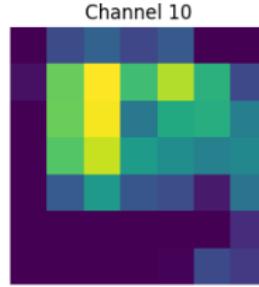
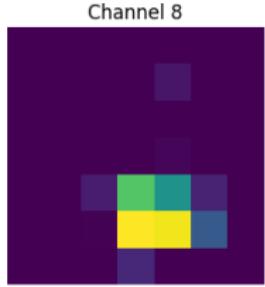
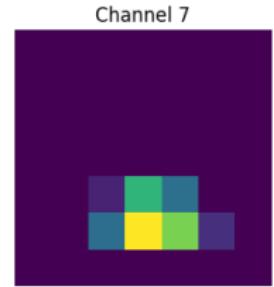
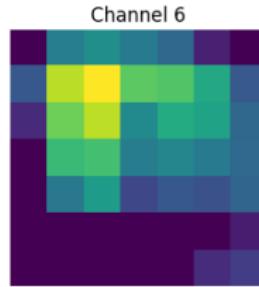
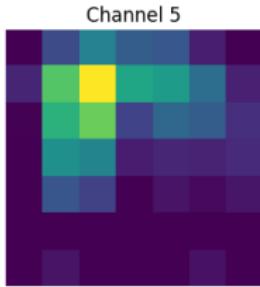
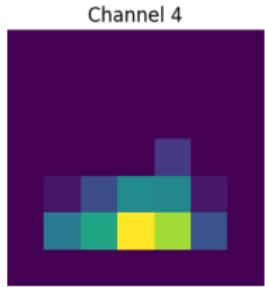
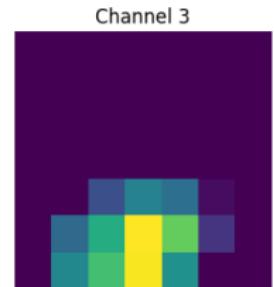
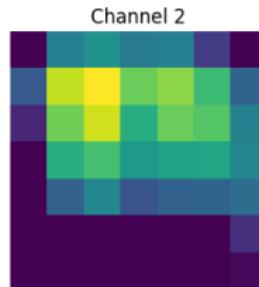
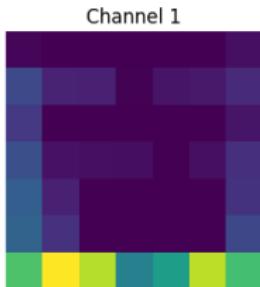
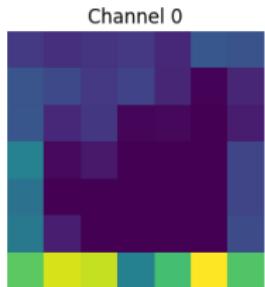
Channel 14

Channel 15

## Feature Maps for MobileNetV2 - Layer: features.13.conv.3



## Feature Maps for MobileNetV2 - Layer: features.18.2



Channel 12

Channel 13

Channel 14



Channel 15



# Methods: Knowledge Distillation

**Knowledge Distillation (KD)** trains a smaller model (student) using soft targets from a larger teacher.

- Teacher outputs softened probabilities:

$$p_t = \text{softmax}\left(\frac{z_t}{T}\right) \quad p_s = \text{softmax}\left(\frac{z_s}{T}\right)$$

- Student learns to match the teacher's softened output distribution and also by true labels:

$$\mathcal{L} = \alpha T^2 \text{CE}(p_t, p_s) + (1 - \alpha) \text{CE}(y, \text{softmax}(z_s))$$

- CE: cross-entropy loss
- $T$ : temperature (softens teacher logits)
- $\alpha$ : balance between KD loss and true-label loss
- $y$ : ground-truth one-hot label

# Implementation

## Datasets:

- Unlabeled training data: 10 000 , 6000, 3000
- Labeled data:
  - 420 for fine-tuning,
  - 180 for validation.
- Test set: 5 000

## Environment

- Google Colab (free tier), training primarily on GPU

## Libraries

- `torchvision.models`: MobileNetV2 ( 3M para), ResNet50 (IMAGENET-1K, 25M para)
- `lightly`: SimCLR augmentations and NT-Xent loss
- `torch.nn.functional`: CE, softmax, log\_softmax, kl\_div for KD loss

# Results: SimCLR

Table: SimCLR Test Performance

Train Size	Acc	Prec-C	Rec-C	Prec-D	Rec-D
10k	0.78	0.78	0.78	0.78	0.79
6k	0.76	0.79	0.71	0.73	0.81
3k	0.72	0.72	0.71	0.72	0.72

## Observations

- Modest acc decreases as data shrinks ( $0.78 \rightarrow 0.72$ ).
- Precision/recall remain well-balanced across Cat/Dog.
- Strong stability indicates effective self-supervised features.

# Results: Knowledge Distillation

Table: KD Test Performance

Train Size	Acc	Prec-C	Rec-C	Prec-D	Rec-D
10k	0.92	0.94	0.90	0.90	0.94
6k	0.84	0.93	0.73	0.78	0.94
3k	0.75	0.79	0.69	0.72	0.81

## Observations

- Very strong performance at 10k ( $\text{Acc} = 0.92$ ).
- At 6k, Unbalances among classes as well as precisions and recalls — likely insufficient epochs.
- At 3k, accuracy is 0.75 with reasonable class balance.
- KD preserves useful boundaries even with limited data.

# Comparison & Training Time

## SimCLR vs KD

- KD outperforms SimCLR across all dataset sizes.
- Largest gap at 10k: KD 0.92 vs. SimCLR 0.78.
- SimCLR degrades more gradually + stable balance → more robust to data reduction.
- KD offers higher absolute accuracy.

## Training Time (Epochs)

Train Size	SimCLR	KD
10k	62	46
6k	57	35
3k	59	25

## Insights

- SimCLR generally requires more epochs, slow convergence on small data.
- KD converges faster because the student leverages teacher guidance.
- Efficiency in terms of training time of KD increases as data size decreases.

# Conclusion

## Key Findings

- Lightweight CNNs (MobileNetV2) perform effectively in low-resource settings.
- **Knowledge Distillation** consistently achieves higher accuracy across all dataset sizes.
- **SimCLR** provides stable and robust performance, even with very small labeled datasets.
- KD converges faster, while SimCLR may require more epochs even for small datasets.

## Overall

- Both techniques significantly improve MobileNetV2 under limited data and computation.
- Demonstrates the usefulness of combining CNN efficiency with modern training strategies.

# Future Experiments

- **Hyperparameter Tuning:** Improve performance through better augmentations, learning rates, KD temperature, and hard/soft loss ratios.
- **Model Variations:** Explore different teacher–student pairs and architectures to enhance distillation effectiveness.
- **Richer Distillation:** Transfer intermediate features or attention maps—not just logits—to strengthen student learning.
- **Adaptivity:** Investigate how students learn new classes, avoid forgetting, and generalize to new domains, also limits of their learning capabilities.